



GFS with Dynamic Replication

Team 37:

Archit Pethani (2022101098)

Neel Amrutia (2022101117)





Objective:

Distributed systems require a storage solution that ensures **resilience**, **scalability**, and **high throughput** for data-intensive applications.



Key Challenge:

How to efficiently manage data replication in distributed systems to optimize **performance**, **fault tolerance**, and **resource utilization**?



Solution:

Develop a Google File System (GFS)-inspired distributed file system with a **dynamic replication feature** that adjusts the number of replicas based on real-time metrics like access patterns and server load.

Project Scope

Our goal was to build a GFS-inspired file system with dynamic replication capabilities.

This system will adapt the number of data chunk replicas based on usage patterns.

We aim to increase replicas for high-demand chunks, ensuring data availability.

Seldom-accessed data will see reduced replication, optimizing resource allocation.

This dual approach enhances overall data management efficiency.



Core Architecture

- Master Server: Responsible for managing metadata, chunk distribution, and overseeing initial replication.
 - Chunk Servers: Store file data chunks, manage read/write operations, and execute replication adjustments as instructed by the master.
 - Client Interface: Provides basic file operations (e.g., create, read, write) and collects access metrics to facilitate replication management.
-

Dynamic Replication Design



DEFINING REPLICATION
POLICIES



MONITORING LOAD FOR
CHUNKS AND CHUNK
SERVERS



DYNAMIC REPLICATION
CONTROL

Dynamic Replication Implementation

- Access counts for each chunk are tracked over fixed intervals (here, 15 seconds)
- **Increase Replicas:** When access exceeds an upper threshold, new replicas are created on underloaded chunkservers.
- **Decrease Replicas:** When access falls below a lower threshold, excess replicas are deleted to save resources.
- The master server assigns new replicas to chunkservers with the lowest load, ensuring even distribution.
- Similarly when the number of replicas have to be decreased then chunk is removed from the chunkservers with higher load.

Current Limitations

- **Delayed Adjustments:** Replica adjustments are based on periodic checks, which may introduce lag in response to rapid changes in load.
- **Minimum Replication Constraint:** A hard limit on the minimum number of replicas (e.g., 2) ensures fault tolerance but may hinder optimal resource usage.
- **Dependency on Load Metrics:** Inaccurate load metrics could lead to suboptimal replication decisions



Key Outcomes

- Efficient use of system resources.

- Improved data availability and fault tolerance.

- Enhanced scalability and adaptability to workload variations.



Conclusion

Dynamic replication addresses the challenges of managing distributed storage in data-intensive systems.

The project provides a scalable, resilient, and high-performance solution inspired by the success of GFS.
